

# Hardware-Assisted Virtualization of Neural Processing Units for Cloud Platforms

Yuqi Xue, Yiqi Liu, Lifeng Nai, Jian Huang  
University of Illinois Urbana-Champaign, Google


PRESENTER:

SUSHIL RAJ REGMI

UNIVERSITY OF NORTH TEXAS

# Neural Processing Units (NPUs)

- NPUs are hardware accelerators designed to accelerate neural network computations (matrix multiplication and convolution)
  - Matrix engines (MEs)
    - matrix based computation, matrix multiplication, convolution
  - Vector engines (VEs)
    - element-wise operations, activation functions



# Virtualizing NPUs in cloud

- **Efficient Resource Utilization**

Prevents underutilization of physical NPUs by sharing resources across multiple workloads

Dynamically allocates idle compute units to active workloads

- **Cost efficiency**

Enables a "pay-as-you-go" model

- **Scalability and Flexibility**

Allows dynamic scaling of NPU resources based on real-time workload demands

- **Supports Cloud-Native AI workloads**

Facilitates multi-tenancy and seamless execution of diverse AI workloads

- **Diverse demands on MEs and VEs**

Accommodate diverse work-load demands

# Challenges in virtualizing NPUs

- **Lack of system abstraction support for NPUs**

Existing cloud platforms expose homogenous NPU cores to the user

Need of a flexible system abstraction that allows users to specify the ME/VE resource and dynamic resource scheduling

- **Lack of architectural support for NPU virtualization**

Existing NPU sharing approach sacrifice isolation or suffer from high preemption overload

Need architectural support to achieve both improved performance isolation and NPU utilization.

- **Lack of ISA support for virtualized NPUs**

Number of compute units need to be explicitly specified at the compilation time which limits the NPU utilization at runtime

Need to design NPU ISA to facilitate dynamic resource scheduling

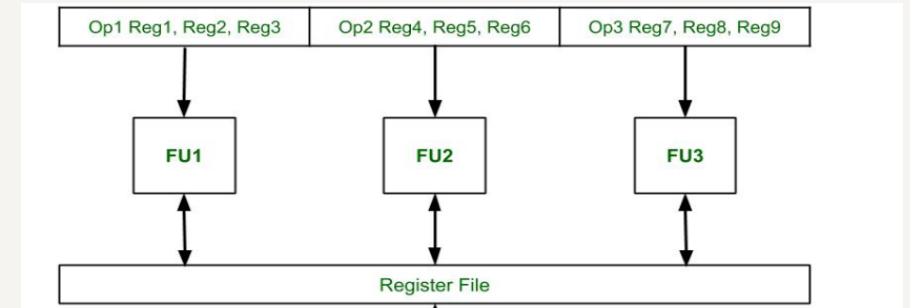
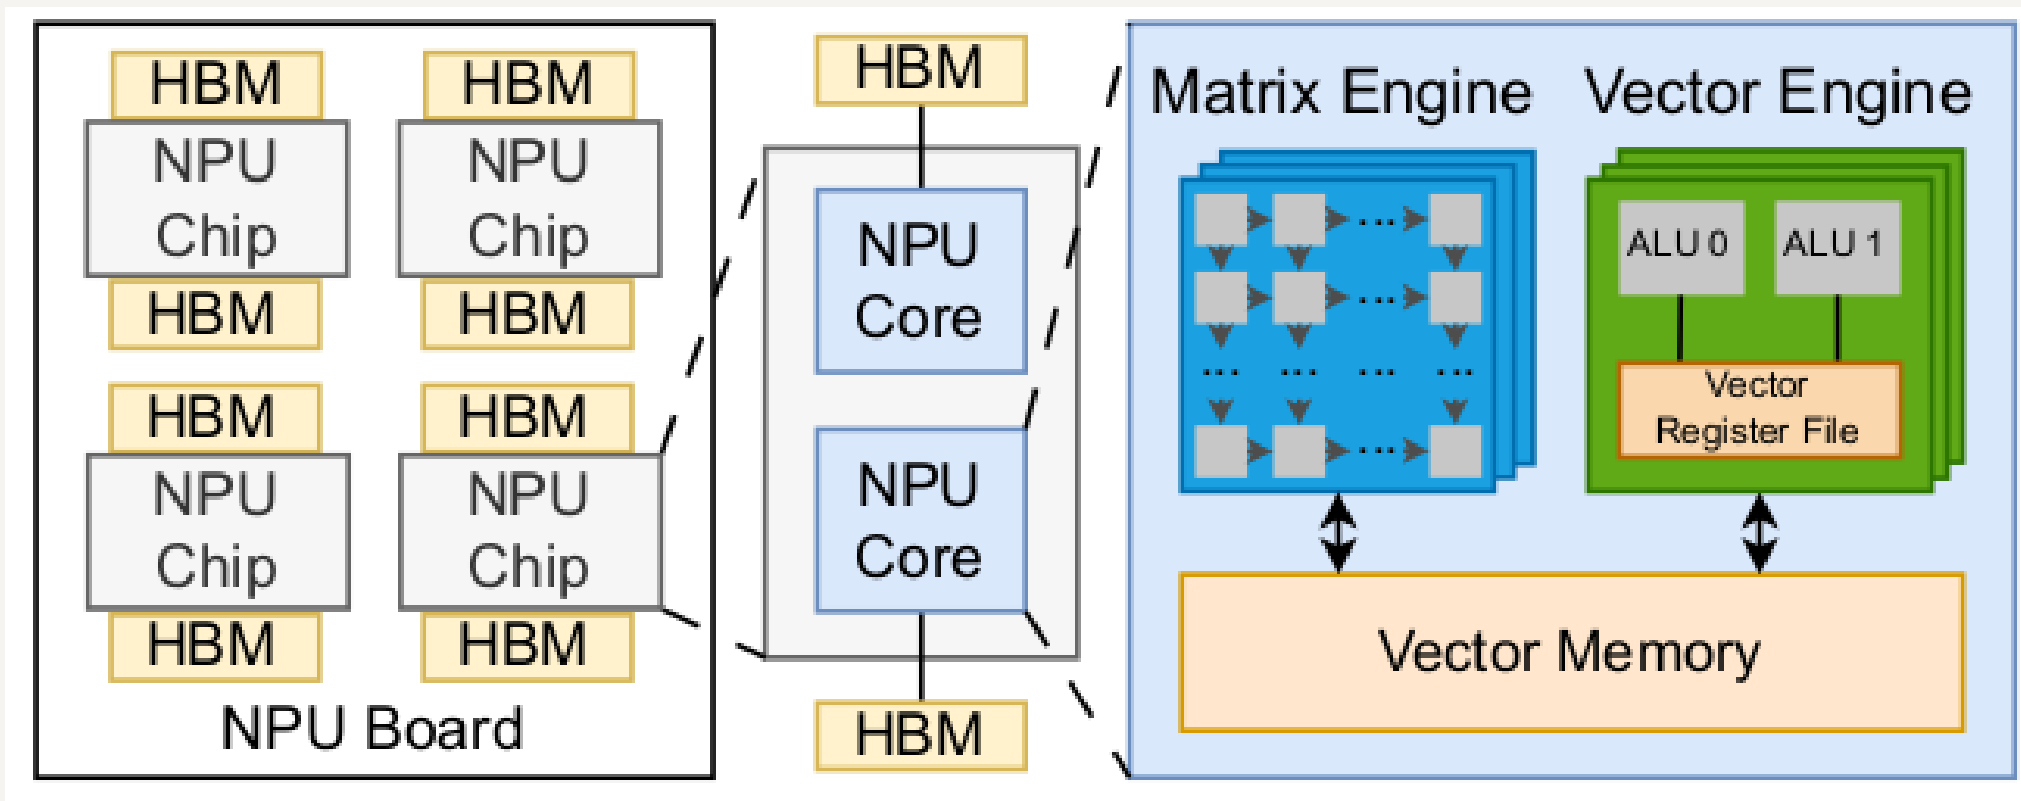



Fig: VLIW style ISA

# NPU System Architecture



# Characterization of ML inference services

Diverse demands on MEs/VEs

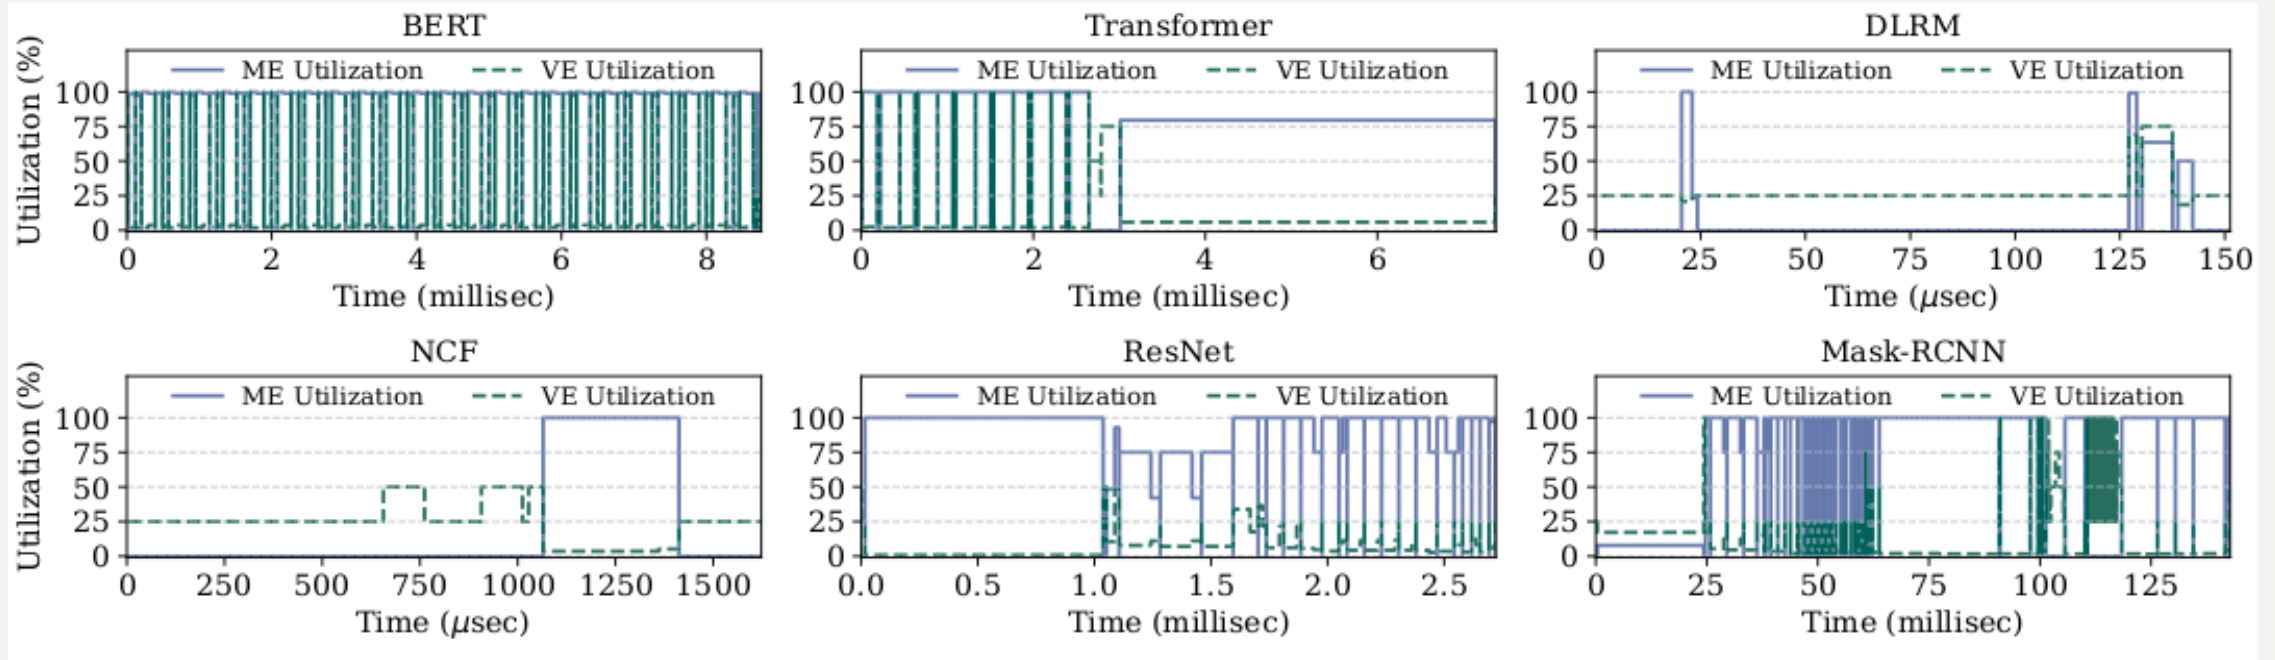


Imbalanced demands are determined by ML model architecture

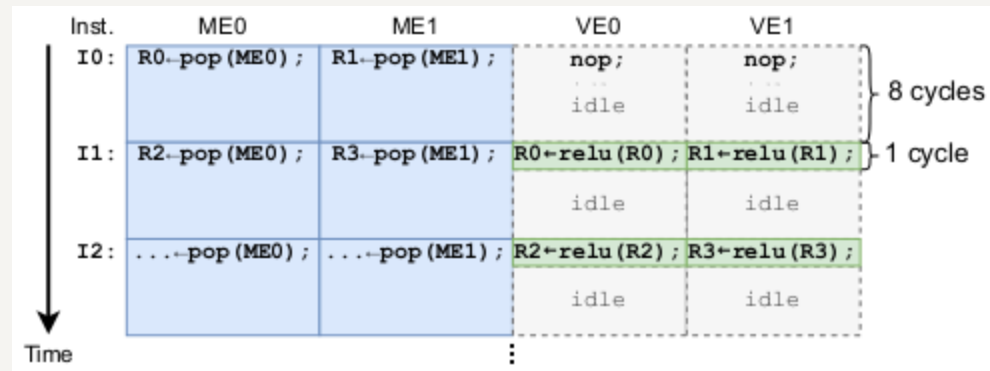
ResNet is dominated by convolutions while DLRM contains many vector operators

# Characterization of ML inference services

## Low NPU resource utilization



## VE underutilization in ME-intensive operator



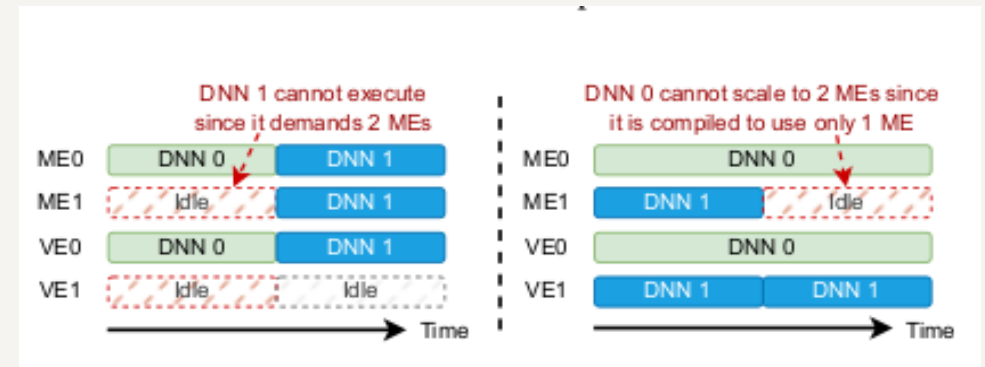
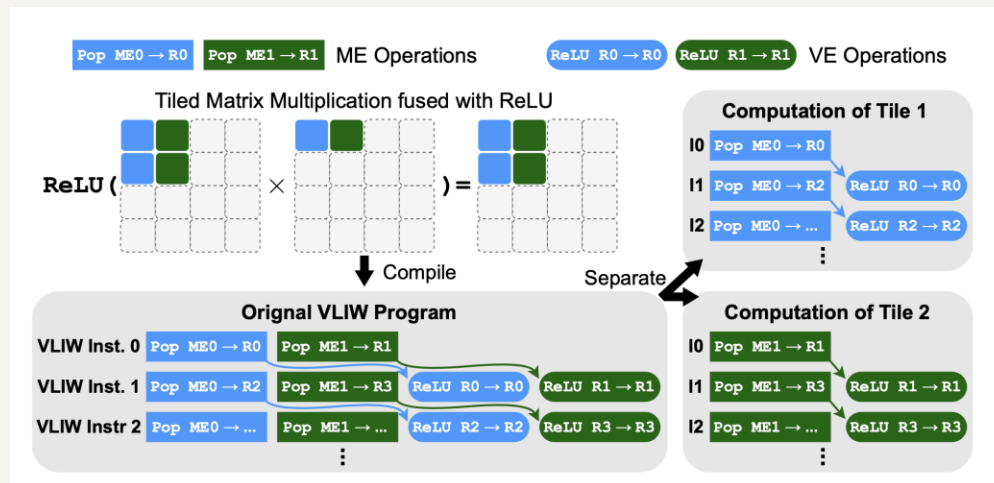
# Opportunities in NPU virtualization

## New abstraction required for fine-grained virtualization

- Need to provide flexibility to users to specify the number of MEs and VEs
- Need to enable dynamic resource scheduling as there is diverse resource demands of different operators over time.

## ISA limitations for enabling virtualized NPU scheduling

- Statically scheduled ISA cannot fully exploit hardware resource at runtime.
- Unnecessarily couple the control flows of all MEs in tensor operator





# Implementation: Neu10

## vNPU: abstraction for NPU virtualization

```
struct vNPU_Config {  
    size_t num_chips;           size_t num_cores_per_chip;  
    size_t num_MEs_per_core;    size_t num_VEs_per_core;  
    size_t sram_size_per_core;  size_t mem_size_per_core;  
};
```

Fig. 10: vNPU configuration.

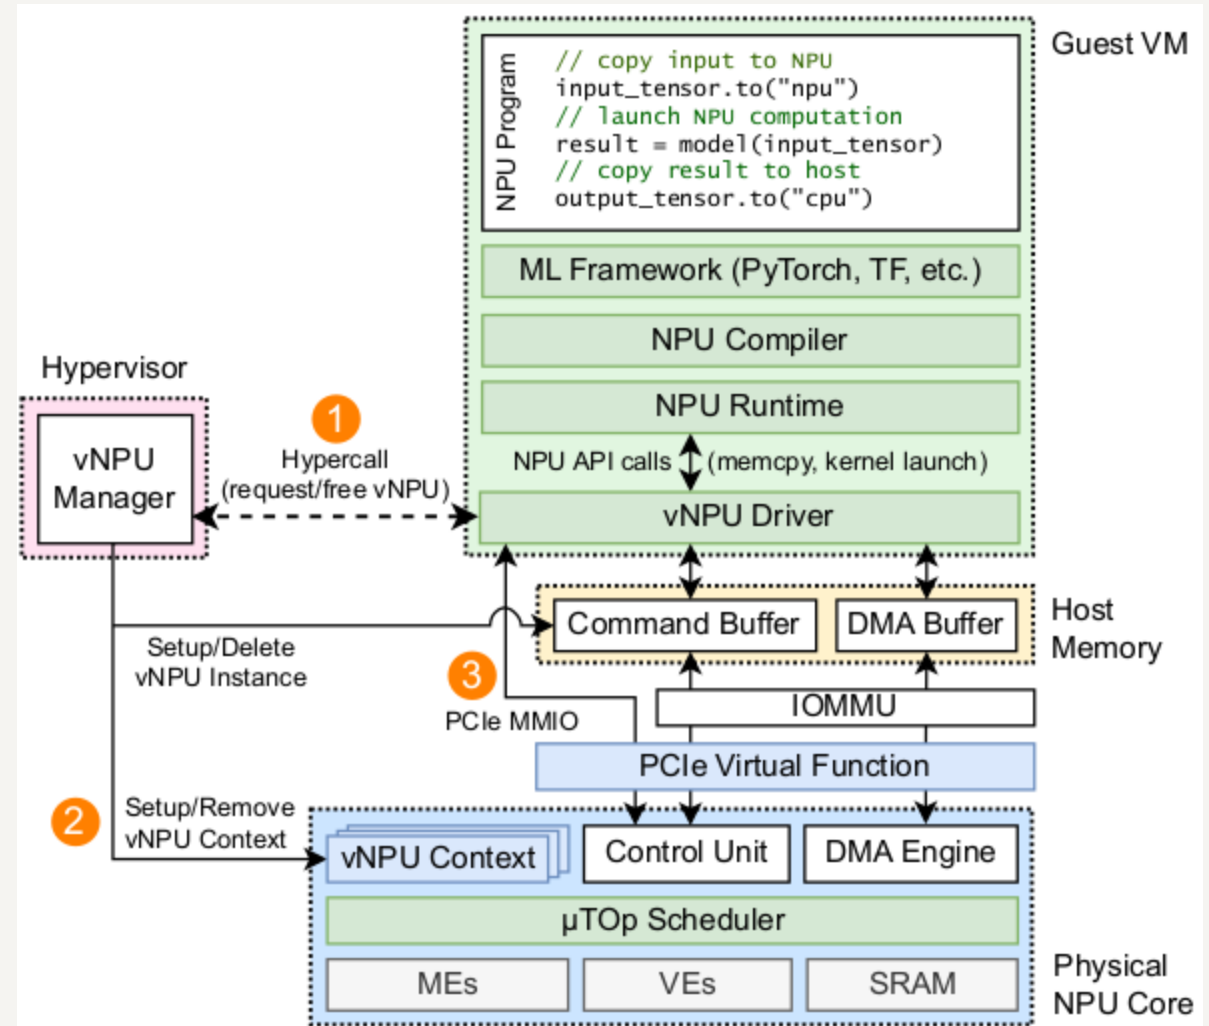


Fig. 11: System architecture of Neu10.

# Neu10: vNPU Allocation and Deallocation

Neu10 allows users to specify the total number of execution units(EUs)

## **ME/VE allocation for workload**

$m$  = ME active runtime/NPU total runtime,  $v$  = VE active runtime/NPU total runtime

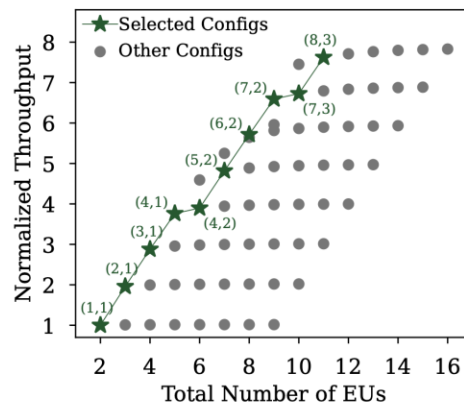
When  $m < 0.5$ , allocate  $\sqrt{m/(1-m)}$  time more MEs than VEs

When  $v < 0.5$ , allocate  $\sqrt{(1-v)/v}$  time more MEs than VEs


When  $m > 0.5$  and  $v > 0.5$ , allocate the same number of VEs and MEs

# Neu10: vNPU Allocation and Deallocation

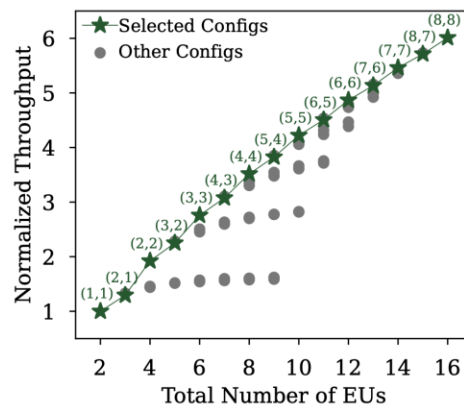
## Cost-effective analysis




(a) BERT (batch size 32).



(b) ResNet (batch size 32).



(c) EfficientNet (batch size 32).



(d) ShapeMask (batch size 8).

Neu10 allocation algorithm selects configuration with better performance than others for same number of EUs

# Neul0: vNPU mapping

vNPUs with many EUs and small memory collocated with vNPUs with few EUs and large memory

vNPU mapping scheme:

- Hardware isolated (spatial isolated) mapping
  - vNPU is mapped to dedicated EUs and SRAM
  - Collocates a set of vNPUs as long as the total resource requirement does not exceed the pNPU.
- Software isolated (temporal sharing) mapping
  - Multiple vNPUs temporally share the same EUs
  - Aims to load-balance the pNPUs while allowing oversubscription

# NeuIO: ISA Extension for NPU virtualization

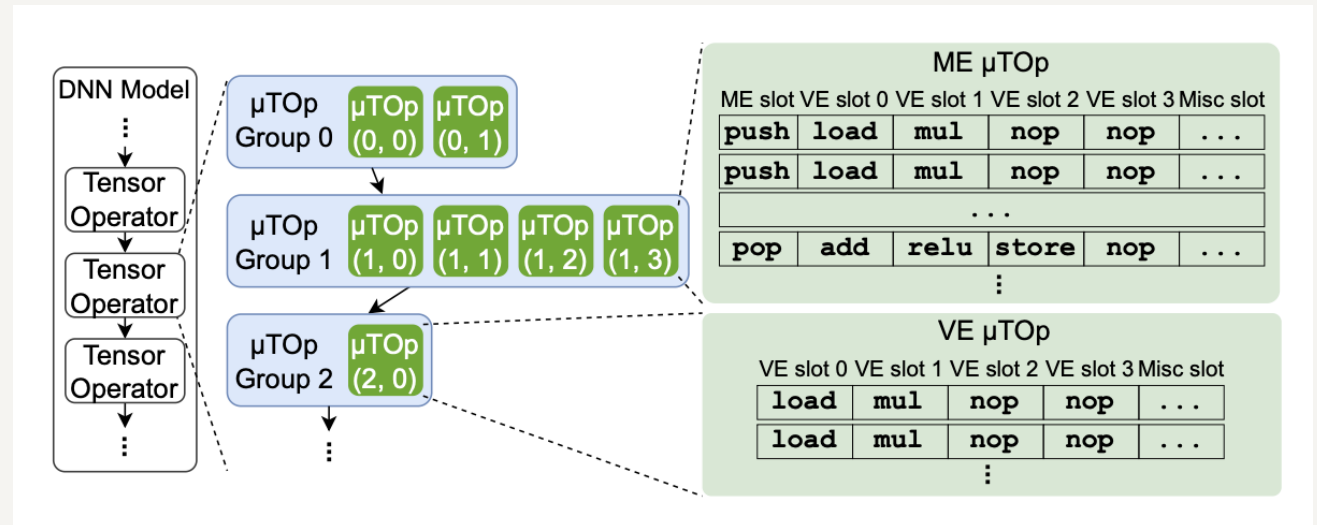
**NeuISA:** Enables dynamic ME/VE scheduling by decomposing tensor operators into runtime-scheduled "sub-tasks"

**$\mu$ TOPs in NeuISA:** Decouples ME/VE control flows into independent instruction sequences ( $\mu$ TOPs)

## $\mu$ TOPs Types:

For pNPU with  $n_x$  MEs and  $n_y$  VEs

- ME  $\mu$ TOP: contains instruction with one ME slot and  $n_y$  VE slots
- VE  $\mu$ TOP: contains instruction with no ME slot and  $n_y$  VE slots



Support fused operator ( ReLU x MatMul ) by organizing  $\mu$ TOPs into a sequence of  $\mu$ TOPs groups.

# Neu0: Architectural support for NeuISA

Hardware scheduler for NeuISA

**Dynamic  $\mu$ TOP Selection:**  $\mu$ TOP scheduler selects the  $\mu$ TOPs to be executed next

**Operation Scheduler:** selects which operation from the instruction queues to be executed at every cycle.

**Low-Overhead Preemption:** Saves ME/VE states to SRAM for fast context switching when reclaiming a harvested ME.

**Parallel Queues:** Dedicated ME/VE queues ( $n_x + n_y$ ) enable concurrent execution.

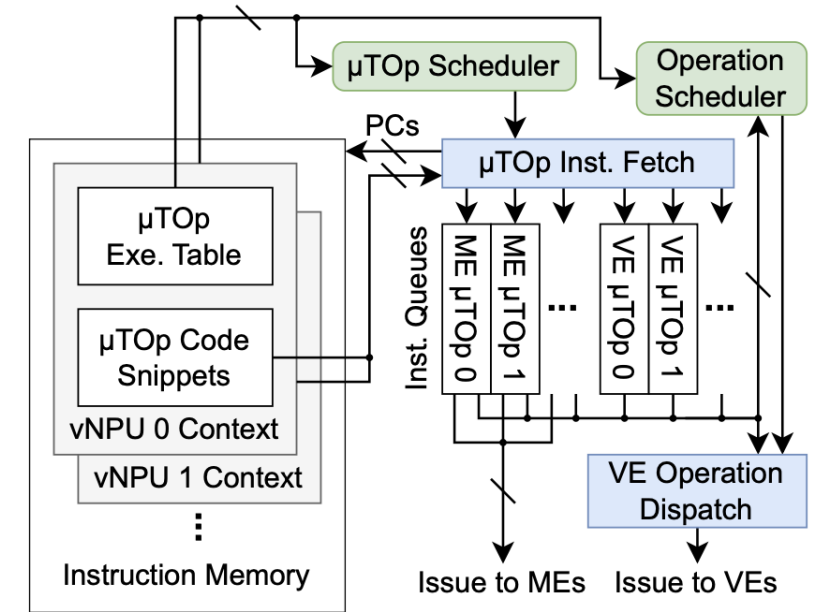




Fig. 17: NPU core pipeline frontend for NeuISA.



# Neu0: Implementation

Implemented Neu10 with production level event-driven NPU simulator.

Obtained operator execution traces (ME/VE time, HBM time, tensor shapes, tile size) on real Google Cloud TPUs

Modify the frontend of simulator to implement the scheduling and harvesting policy

Prototyped the hardware scheduler in Verilog

Hardware area overhead of Neu10 is only 0.04% on a TPUv4 chip

# of MEs/VEs	4 MEs & 4 VEs
ME dimension	$128 \times 128$ systolic array
VE ALU dimension	$128 \times 8$ FP32 operations/cycle
Frequency	1050 MHz
On-chip SRAM	128 MB
HBM Capacity & Bandwidth	64 GB, 1200 GB/s

# Evaluation

## Setup:

### Workload combination

- **Low ME/VE contention:** (DLRM+SMask, DLRM+RtNT, NCF+RsNt)
- **Medium ME/VE contention:** (ENet+SMask, BERT+ENet, ENet+MRCN)
- **High ME/VE contention:** (ENet+TFMR, MNIST+RtNt, RNRS+RtNt)

Each workload runs on a vNPU with 2 MEs and 2 VEs

Maps 2 vNPUs to physical NPU core with 4 MEs and 4 VEs

Benchmark for comparison:

**PMT:** temporal-sharing of the entire NPU core among multiple vNPUs.

**V10:** temporal-sharing of all MEs and VEs among the vNPUs with priority based preempting policy

**Neu10-NoHarvest (Neu10-NH):** spatial-isolated vNPUs with dedicated MEs/VEs without dynamic scheduling.

Category	Model Name	Abbrev.	HBM Footprint (batch size = 8)
Natural Language Processing	BERT	BERT	1.27GB
	Transformer	TFMR	1.54GB
Recommendation	DLRM	DLRM	22.38GB
	NCF	NCF	11.10GB
Object Detection & Segmentation	Mask-RCNN	MRCNN	3.21GB
	RetinaNet	RtNt	860.51MB
	ShapeMask	SMask	6.04GB
Image Classification	MNIST	MNIST	10.59MB
	ResNet	RsNt	216.02MB
	ResNet-RS	RNRS	458.17MB
	EfficientNet	ENet	99.06MB



# Evaluation

## Tail Latency:

Neu10 improves the 95% tail latency over V10 by upto 1.56x on average




Fig. 19: 95% Percentile latency of Neu10 (normalized to PMT).

## Average Latency:

Neu10 improves the average latency by 1.33x over PMT and 1.12x over V10




Fig. 20: Average request latency of Neu10 (normalized to PMT).

## Throughput:

- V10 and Neu10 improve the throughput over PMT (by  $1.58\times$  and  $1.62\times$  on average), by overlapping the execution of ME intensive operators and VE-intensive operators.
- When the ME/VE contention is high, Neu10 improves the throughput of DNN workloads over V10 by up to  $1.41\times$ .




Fig. 21: Throughput of Neu10 (normalized to PMT).


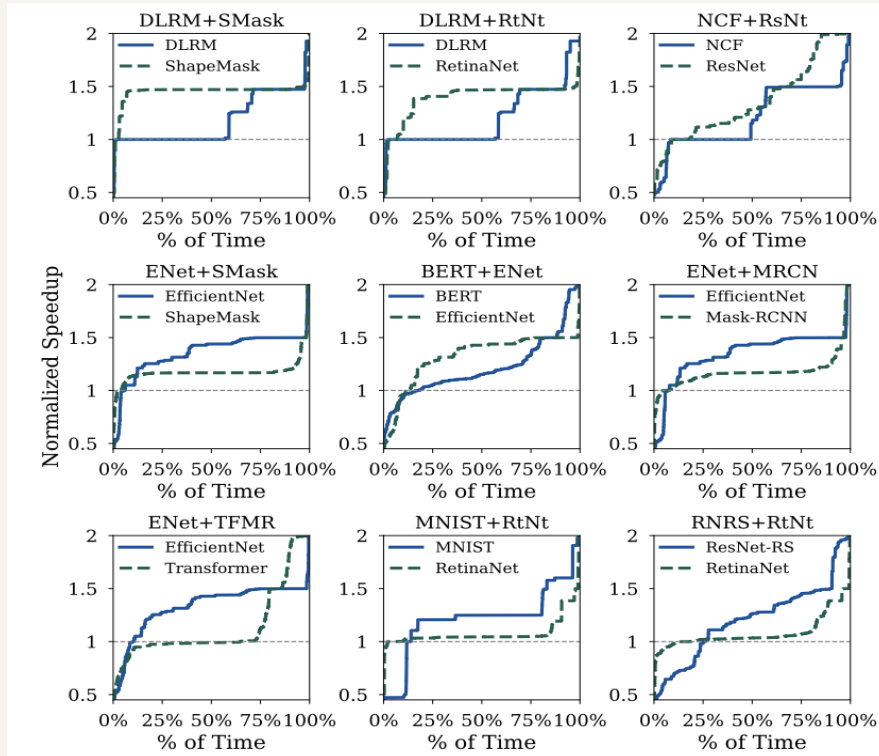
# Evaluation

## Resource Utilization Improvement:


1.26x ME and 1.2x VE utilization improvement over PMT on average

## Benefit of ME/VE harvesting

Speed up of each operator in Neu10 over Neu10-NH



(a) Total ME utilization of the NPU core.



(b) Total VE utilization of the NPU core.

# Evaluation

## Impact of varying MEs and VEs

Neu10 benefits since there is more flexibility for dynamic ME/VE scheduling




Fig. 25: Throughput improvement of Neu10 with varying numbers of MEs and VEs over V10 with 2 MEs and 2 VEs.

## Impact of varying Memory Bandwidth

Comparable throughput

For memory intensive workloads, Neu10 brings more benefit with more bandwidth




Fig. 26: Throughput improvement of Neu10 with varying HBM bandwidth (normalized to V10).

# Conclusion

- Highlighted challenges in NPU virtualization, including lack of fine-grained abstraction, rigid scheduling, and architectural limitations.
- Proposed Neu10, a framework enabling fine-grained abstraction, dynamic resource allocation, and hardware-assisted virtualization.
- Demonstrated improved utilization, reduced tail latency, and better performance isolation for multi-tenant ML workloads.
- Paved the way for scalable, efficient, and flexible NPU resource management in cloud platforms to meet growing AI demands.